

# CS-470: Machine Learning

## Week 11 - Training Neural Networks: A Deep Dive

Instructor: Dr. Sajjad Hussain

Department of Electrical and Computer Engineering  
SEECS, NUST

November 27th, 2025

# Overview

- 1 The Problem: Vanishing Gradients
- 2 Smart Initialization
- 3 Beyond ReLU: Other Activation Functions
- 4 Batch Normalization
- 5 Transfer Learning
- 6 Optimizers
- 7 Regularization

# The Vanishing Gradient Problem

- In deep networks, gradients become extremely small as they are backpropagated from the output layer to the initial layers.
- This means **early layers learn very slowly** or not at all, while later layers converge.
- The core of the problem lies in the **chain rule** and the **activation function**.

## The Chain Rule in Backpropagation

For a weight  $w^{(1)}$  in the first layer, the gradient is:

$$\frac{\partial \mathcal{L}}{\partial w^{(1)}} = \underbrace{\frac{\partial \mathcal{L}}{\partial a^{(L)}}}_{\text{Output Grad}} \cdot \underbrace{\frac{\partial a^{(L)}}{\partial a^{(L-1)}} \cdots \frac{\partial a^{(2)}}{\partial a^{(1)}}}_{\text{Many terms}} \cdot \frac{\partial a^{(1)}}{\partial z^{(1)}} \cdot \frac{\partial z^{(1)}}{\partial w^{(1)}}$$

Each term  $\frac{\partial a^{(l)}}{\partial a^{(l-1)}}$  depends on the derivative of the activation function,  $g'(z^{(l-1)})$ .


# Why Sigmoid Causes Vanishing Gradients

The Sigmoid function and its derivative:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

$$\sigma'(x) = \sigma(x)(1 - \sigma(x))$$

- Maximum value of  $\sigma'(x)$  is 0.25.
- For most inputs ( $|x| > 4$ ),  
 $\sigma'(x) \approx 0$ .



**The Problem:**  $g'(z) \leq 0.25$ . In a deep network, multiplying many such small numbers ( $< 1$ ) causes the product to **vanish exponentially**.

$$\frac{\partial \mathcal{L}}{\partial w^{(1)}} \propto (0.25 \times 0.25 \times \dots \times 0.25) \approx 0$$


# The ReLU Solution

The ReLU (Rectified Linear Unit) function:

$$\text{ReLU}(x) = \max(0, x)$$

Its derivative:

$$\text{ReLU}'(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$



## How ReLU Helps:

- For active neurons ( $x > 0$ ), the derivative is **exactly 1**.
- This prevents the gradient from shrinking *purely* due to the activation function when the neuron is active.
- The product in the chain rule is now  $\propto 1 \times 1 \times \cdots \times 1 = 1$ , preventing exponential vanishing.

# Glorot (Xavier) and He Initialization

**The Goal:** Prevent the outputs of the layers from exploding or vanishing *during the forward pass*, even in deep networks.

## Intuition

We want the **variance** of the outputs of each layer to be equal to the variance of its inputs. Similarly, we want the gradients to have the same variance during backpropagation.

For a linear layer:  $z = Wx + b$ . Assuming  $E[W] = 0$ ,  $E[x] = 0$ :

$$\text{Var}(z) = n_{\text{in}} \cdot \text{Var}(W) \cdot \text{Var}(x)$$

To have  $\text{Var}(z) = \text{Var}(x)$ , we need:

$$n_{\text{in}} \cdot \text{Var}(W) = 1 \quad \Rightarrow \quad \text{Var}(W) = \frac{1}{n_{\text{in}}}$$

Considering backpropagation as well, Glorot et al. proposed a compromise.

# Glorot vs. He Initialization

## Glorot/Xavier Initialization (for Tanh/Sigmoid)

Balances the variance for both forward and backward passes.

$$W \sim \mathcal{N} \left( 0, \sqrt{\frac{2}{n_{\text{in}} + n_{\text{out}}}} \right) \quad \text{or} \quad \mathcal{U} \left( -\sqrt{\frac{6}{n_{\text{in}} + n_{\text{out}}}}, \sqrt{\frac{6}{n_{\text{in}} + n_{\text{out}}}} \right)$$

## He Initialization (for ReLU and variants)


Accounts for the fact that ReLU sets half its inputs to zero, effectively halving the variance. It focuses on the forward pass variance.

$$W \sim \mathcal{N} \left( 0, \sqrt{\frac{2}{n_{\text{in}}}} \right) \quad \text{or} \quad \mathcal{U} \left( -\sqrt{\frac{6}{n_{\text{in}}}}, \sqrt{\frac{6}{n_{\text{in}}}} \right)$$

**Summary:** Use the right initialization for your activation function to get stable training from the start.

# Variants of ReLU


## Leaky ReLU



$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha x & \text{if } x \leq 0 \end{cases}$$

Prevents "dying ReLU".  $\alpha$  is a small constant (e.g., 0.01).


## Parametric ReLU (PReLU)



$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha x & \text{if } x \leq 0 \end{cases}$$

The slope  $\alpha$  is a **learnable parameter**.

## Exponential LU (ELU)



$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha(e^x - 1) & \text{if } x \leq 0 \end{cases}$$

Smooth transition, helps push mean activations towards zero.


# Scaled Exponential LU (SELU)

**SELU** is a self-normalizing activation function.

$$\text{SELU}(x) = \lambda \begin{cases} x & \text{if } x > 0 \\ \alpha(e^x - 1) & \text{if } x \leq 0 \end{cases}$$

With  $\alpha \approx 1.6733$  and  $\lambda \approx 1.0507$ .

- Under specific conditions (LeCun initialization), SELU networks **self-normalize**: the outputs of each layer tend to preserve zero mean and unit variance.
- This inherently avoids vanishing/exploding gradients.
- Very effective for deep feedforward networks.



# Batch Normalization: The Idea

**Internal Covariate Shift:** The change in the distribution of network activations due to the change in network parameters during training. This slows down training.

## BatchNorm Solution

Normalize the inputs of each layer to have **zero mean** and **unit variance**, for each mini-batch. This stabilizes the distribution of activations.



**Figure:** BatchNorm is applied *after* the linear transformation and *before* the non-linear activation.

# Batch Normalization: The Algorithm

For a layer output (pre-activation)  $\mathbf{z}$  over a mini-batch  $\mathcal{B} = \{z_1, \dots, z_m\}$ :

- ➊ Calculate the mean and variance of the batch:

$$\mu_{\mathcal{B}} = \frac{1}{m} \sum_{i=1}^m z_i \quad \sigma_{\mathcal{B}}^2 = \frac{1}{m} \sum_{i=1}^m (z_i - \mu_{\mathcal{B}})^2$$

- ➋ Normalize the values:

$$\hat{z}_i = \frac{z_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}}$$

( $\epsilon$  is a small constant for numerical stability)

- ➌ Scale and shift (this is crucial!):

$$\text{BN}(z_i) = \gamma \hat{z}_i + \beta$$

**Key:**  $\gamma$  (scale) and  $\beta$  (shift) are **learnable parameters**. They allow the network to *decide* if normalization is useful, and to represent the identity function if that's optimal.

# Why BatchNorm Works So Well

- **Allows Higher Learning Rates:** By stabilizing gradients, it prevents updates from blowing up the model.
- **Reduces Sensitivity to Initialization:** Less careful initialization is needed.
- **Acts as a Regularizer:** The output for a given example depends on the statistics of the entire mini-batch, adding noise. This reduces overfitting.
- At test time, we use **running averages** of the mean and variance computed during training, not the batch statistics.

# Transfer Learning: Leveraging Pre-trained Models

**"Why train from scratch when someone has already done the hard work?"**

## Scenarios:

- ① **Small target dataset, similar to source:** Freeze all convolutional layers, train only the classifier head.
- ② **Medium target dataset, similar to source:** Freeze early layers, fine-tune later layers and the classifier.
- ③ **Large target dataset, similar to source:** Fine-tune the entire network (this is just using the pre-trained weights as a very good initialization).
- ④ **Dataset not similar to source:** Transfer learning might not help much. Consider training from scratch.

# Transfer Learning




Figure: Transfer Learning<sup>1</sup>

<sup>1</sup>Image adapted from Aurélien Géron, "Hands-On Machine Learning with Scikit-Learn, Keras & TensorFlow"

# Optimizers: Vanilla Stochastic Gradient Descent (SGD)

## The Core Update Rule

The simplest form of gradient descent updates the parameters  $\theta$  using only the current gradient  $\nabla J(\theta_t)$ :

$$\theta_{t+1} = \theta_t - \eta g_t \nabla_{\theta} J(\theta_t)$$

where  $g_t = \nabla_{\theta} J(\theta_t)$  is the current gradient and  $\eta$  is the learning rate.

## The Problem:

- The update is *myopic*; it only looks at the gradient **at this very moment**.
- As we approach a minimum, the gradient  $\nabla J(\theta_t)$  naturally becomes smaller.
- Therefore, the update step  $-\eta \nabla J(\theta_t)$  also becomes smaller, leading to very slow convergence near the optimum.
- It has no memory of past gradients, making it susceptible to getting stuck in saddle points and sensitive to noisy gradients.

**Summary:** SGD slows down near the minimum because its update is directly proportional to the gradient, which vanishes at the optimum.

# Optimizers: SGD with Momentum

**The Idea:** Simulate a heavy ball rolling downhill. It doesn't just follow the current slope but maintains **momentum** from previous gradients.

## The Algorithm

We introduce a velocity vector  $v_t$  that accumulates past gradients.

$$v_t = \beta v_{t-1} + \eta g_t$$

$$\theta_{t+1} = \theta_t - v_t$$

- $\beta$  is the momentum term (e.g., 0.9). It dictates how much of the past velocity is retained.
- $v_t$  is the **exponentially decaying average** of all past gradients.

**Mathematical Insight:** Let's unroll the velocity term for the first few steps:

$$v_0 = 0, \quad v_1 = \beta v_0 + \eta g_1 = \eta g_1$$

$$v_2 = \beta v_1 + \eta g_2 = \beta \eta g_1 + \eta g_2$$

# SGD with Momentum: Why It Works

$$\begin{aligned}v_3 &= \beta v_2 + \eta g_3 = \beta(\beta\eta g_1 + \eta g_2) + \eta g_3 = \beta^2\eta g_1 + \beta\eta g_2 + \eta g_3 \\v_t &= \eta(g_t + \beta g_{t-1} + \beta^2 g_{t-2} + \dots)\end{aligned}$$

The update is a weighted sum of **all previous gradients**, with more recent gradients weighted more heavily.

## Benefits:

- **Faster Convergence:** In directions with a consistent gradient, momentum adds up, leading to larger updates ( $v_t$  is large).
- **Reduces Oscillations:** In ravines (steep walls, gentle slope), gradients oscillate across the sides. The momentum term cancels out these opposing gradients, leading to a smoother, faster path down the gentle slope.
- **Escapes Flat Regions:** Even if the current gradient is near zero, the accumulated momentum can carry the parameters through.

# Nesterov Accelerated Gradient (NAG)

**"Look-ahead" Momentum:** A smarter version of momentum.

## Nesterov Momentum

- ① First, make a "**look-ahead**" jump using the old velocity:

$$\theta_{\text{look-ahead}} = \theta_t - \beta v_{t-1}.$$

- ② Then, calculate the gradient **at this future position**,  $g_t = \nabla J(\theta_t - \beta v_{t-1})$ .
- ③ Finally, update the velocity and parameters:

$$v_t = \beta v_{t-1} + \eta g_t$$

$$\theta_{t+1} = \theta_t - v_t$$

**Why it's Better:** By calculating the gradient *after* the momentum jump, NAG gets a "preview" of where it's going. This creates a **corrective term**. If the momentum jump was too big and is about to increase the loss, the gradient at the look-ahead point will correct the velocity, preventing overshooting and leading to more responsive behavior.

# RMSProp: Adaptive Learning Rates

**The Idea:** Adjust the learning rate for **each parameter** individually based on the history of its gradients.

## The Algorithm

Accumulate an exponentially decaying average of **squared gradients**.

$$E[g^2]_t = \beta E[g^2]_{t-1} + (1 - \beta)g_t^2$$

$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}} g_t$$

## Effect:

- Parameters with large, frequent gradients  $\Rightarrow$  Large  $E[g^2]_t \Rightarrow$  **Small effective learning rate**  $\frac{\eta}{\sqrt{E[g^2]_t}}$ .
- Parameters with small, infrequent gradients  $\Rightarrow$  Small  $E[g^2]_t \Rightarrow$  **Large effective learning rate**.

This automatically normalizes the step size for each parameter, which is very effective for problems with sparse gradients or ill-conditioned landscapes.

# Adam: Adaptive Moment Estimation

**The Best of Both Worlds:** Adam combines the concepts of **Momentum** (first moment) and **RMSProp** (second moment).

## The Algorithm

Compute gradient:  $g_t = \nabla_{\theta} J(\theta_t)$

Update biased first moment (momentum):  $m_t = \beta_1 m_{t-1} + (1 - \beta_1)g_t$

Update biased second moment (RMS):  $v_t = \beta_2 v_{t-1} + (1 - \beta_2)g_t^2$

Compute bias-corrected first moment:  $\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$

Compute bias-corrected second moment:  $\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$

Update parameters:  $\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} \hat{m}_t$

Typical values:  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$ ,  $\epsilon = 10^{-8}$ .

# Adam: Understanding Bias Correction

- $\hat{m}_t$  provides a smooth, momentum-like direction.
- $\hat{v}_t$  adapts the learning rate for each parameter, like RMSProp.
- **Bias Correction** is crucial in early steps when  $m_t$  and  $v_t$  are initialized to 0.

## What are $\beta_1^t$ and $\beta_2^t$ ?

- They represent  $\beta_1$  and  $\beta_2$  raised to the power of  $t$
- $t$  is the current time step (iteration number)
- Used in the bias correction terms:

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t} \quad \text{and} \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

## Why Bias Correction is Needed

- At  $t = 0$ ,  $m_0 = 0$  and  $v_0 = 0$
- Early estimates are **biased toward zero**
- Example:  $m_1 = 0.9 \cdot 0 + 0.1 \cdot g_1 = 0.1g_1$  (too small!)
- Without correction, early updates would be too small

Adam is robust and works well out-of-the-box on a wide variety of problems.

# L1 and L2 Regularization

**Goal:** Prevent overfitting by penalizing large weights, encouraging a simpler model.

## L2 Regularization (Weight Decay, Ridge)

Adds the squared magnitude of weights to the loss function.

$$J_{\text{reg}}(\theta) = J(\theta) + \frac{\alpha}{2} \|\theta\|_2^2$$

**Effect:** Shrinks all weights proportionally towards zero. The gradient is  $\nabla J(\theta) + \alpha\theta$ .

## L1 Regularization (Lasso)

Adds the absolute magnitude of weights to the loss function.


$$J_{\text{reg}}(\theta) = J(\theta) + \alpha \|\theta\|_1$$

**Effect:** Can drive some weights **exactly to zero**, creating a sparse model. Useful for feature selection.


# Dropout

**Idea:** During training, randomly "drop out" (set to zero) a fraction  $p$  of the neurons in a layer for each forward pass. This prevents complex co-adaptations on training data.

Standard Network



Network with Dropout Applied



**Figure:** Dropout forces the network to learn redundant representations.

# How Dropout Works

- During each training iteration (for each mini-batch), we randomly "**drop**" (set to zero) a fraction  $p$  of the neurons in the dropout-applied layers.
- This creates a different **random sub-network** for each training example, where only a subset of neurons are active and contribute to the forward and backward passes.
- At test time, we use the **full network**. To compensate for the fact that more neurons are active, the weights of the trained network are **scaled down by**  $1 - p$  (or activations are scaled up by  $1/(1 - p)$  during training).
- **Effect:** It's like training a large ensemble of many different sub-networks and averaging their predictions at test time.
- This prevents overfitting very effectively and is a key technique in modern deep learning.

# Summary

- **Vanishing Gradients:** Solved by ReLU and proper initialization (He/Xavier).
- **Initialization:** Critical for stable training. Use He for ReLU, Glorot for Tanh.
- **Activation Functions:** ReLU is standard; Leaky ReLU/ELU/SELU solve the "dying ReLU" problem.
- **BatchNorm:** Normalizes layer inputs, allowing higher learning rates and acting as a regularizer.
- **Transfer Learning:** Use pre-trained models and fine-tune based on your data size and similarity.
- **Optimizers:** Adam is a good default, combining Momentum and RMSProp.
- **Regularization:** L1/L2 penalize weights; Dropout prevents feature co-adaptation.