

# CS-470: Machine Learning

## Week 12 - Clustering Algorithms

Instructor: Dr. Sajjad Hussain

Department of Electrical and Computer Engineering  
SEECS, NUST

December 4th, 2025

# What is Clustering?

## Definition

Clustering is an **unsupervised learning** technique that groups similar data points together while keeping dissimilar points in different groups.

- **Goal:** Discover natural groupings in data
- **Input:** Unlabeled data points
- **Output:** Groups (clusters) where intra-cluster similarity is high and inter-cluster similarity is low

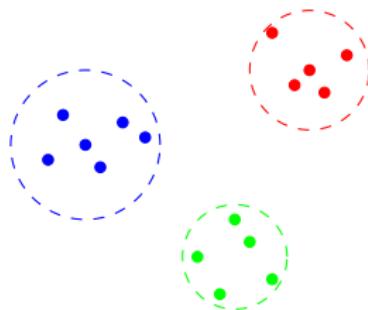


Figure: Example of clustered data with three clusters

# Applications of Clustering

- **Customer Segmentation**
- **Image Segmentation**
- **Anomaly Detection**
- **Document Classification**
- **Social Network Analysis**
- **Market Research**
- **Gene Sequence Analysis**
- **Recommender Systems**

## Real-World Example: Customer Segmentation

- Group customers by purchasing behavior
- Targeted marketing campaigns
- Personalized recommendations
- Customer retention strategies

# K-Means: Basic Idea

## Objective Function

Minimize the within-cluster sum of squares (WCSS):

$$J = \sum_{i=1}^k \sum_{x \in C_i} \|x - \mu_i\|^2$$

where:

- $k$ : number of clusters
- $C_i$ : set of points in cluster  $i$
- $\mu_i$ : centroid of cluster  $i$

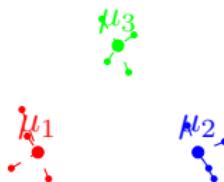


Figure: K-Means minimizes distances from points to their cluster centroids



# K-Means Algorithm

## Algorithm 1 K-Means Clustering

**Require:** Data points  $X$ , number of clusters  $k$ , maximum iterations  $T$

1: Initialize  $k$  centroids  $\mu_1, \mu_2, \dots, \mu_k$

2: **for**  $t = 1$  to  $T$  **do**

3:   **Assignment Step:** Assign each point to nearest centroid

$$C_i = \{x : \|x - \mu_i\|^2 \leq \|x - \mu_j\|^2 \ \forall j\}$$

4:   **Update Step:** Recalculate centroids

$$\mu_i = \frac{1}{N_i} \sum_{x \in C_i} x$$

5:   **if** centroids don't change **then**

6:       **break**

7:   **end if**

8: **end for**

9: **return** Clusters  $C_1, C_2, \dots, C_k$  and centroids  $\mu_1, \mu_2, \dots, \mu_k$



# Toy Example: Step 1 - Initialization

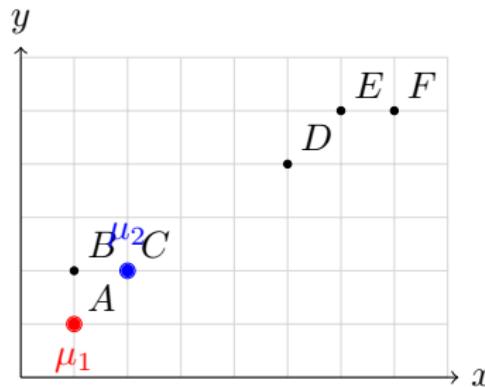
## Data Points

Points:  $A(1, 1)$ ,  $B(1, 2)$ ,  $C(2, 2)$ ,  $D(5, 4)$ ,  $E(6, 5)$ ,  $F(7, 5)$

Choose  $k = 2$  clusters

## Random Initial Centroids

Let's initialize:  $\mu_1 = (1, 1)$  (point A),  $\mu_2 = (2, 2)$  (point C)



# Step 2 - First Assignment

Calculate distances

Point	Distance to $\mu_1$	Distance to $\mu_2$	Cluster
A(1,1)	0	$\sqrt{(1-2)^2 + (1-2)^2} = 1.41$	$C_1$
B(1,2)	1	$\sqrt{(1-2)^2 + (2-2)^2} = 1$	$C_2$
C(2,2)	1.41	0	$C_2$
D(5,4)	5	$\sqrt{(5-2)^2 + (4-2)^2} = 3.61$	$C_2$
E(6,5)	6.4	$\sqrt{(6-2)^2 + (5-2)^2} = 5$	$C_2$
F(7,5)	7.2	$\sqrt{(7-2)^2 + (5-2)^2} = 5.83$	$C_2$

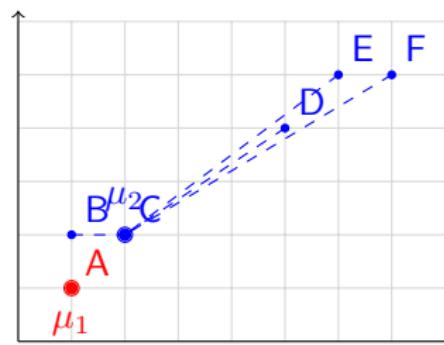


Figure: Assignment: Cluster 1: {A}. Cluster 2: {B,C,D,E,F}

# Step 3 - Update Centroids

## Recalculate Centroids

$$\text{Cluster 1: } \mu_1 = \frac{1}{1}[(1, 1)] = (1, 1)$$

$$\text{Cluster 2: } \mu_2 = \frac{1}{5}[(1, 2) + (2, 2) + (5, 4) + (6, 5) + (7, 5)] = \left(\frac{21}{5}, \frac{18}{5}\right) = (4.2, 3.6)$$

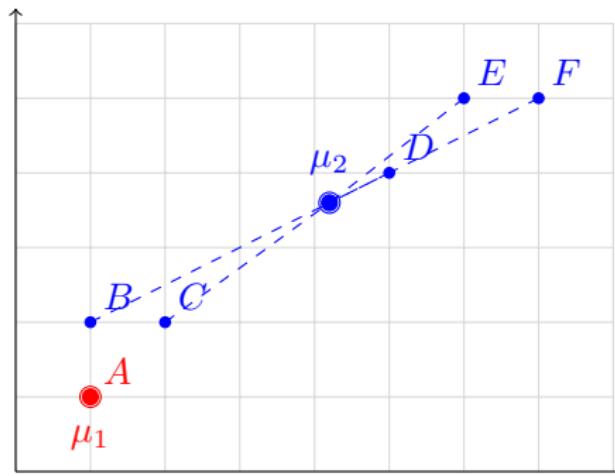


Figure: Updated centroids after first iteration

# Step 4 - Reassignment

Recalculate distances to new centroids

Point	Dist to $\mu_1(1, 1)$	Dist to $\mu_2(4.2, 3.6)$	Cluster
A(1,1)	0	3.83	$C_1$
B(1,2)	1	3.30	$C_1$
C(2,2)	1.41	2.58	$C_1$
D(5,4)	5	1.08	$C_2$
E(6,5)	6.4	2.15	$C_2$
F(7,5)	7.2	3.08	$C_2$

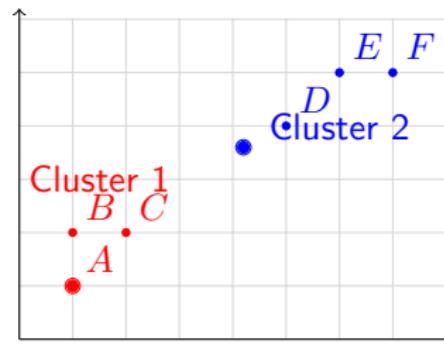


Figure: New assignment: Cluster 1: {A,B,C}, Cluster 2: {D,E,F}

# Step 5 - Final Centroids

## Update Centroids Again

$$\text{Cluster 1: } \mu_1 = \frac{1}{3}[(1, 1) + (1, 2) + (2, 2)] = \left(\frac{4}{3}, \frac{5}{3}\right) = (1.33, 1.67)$$

$$\text{Cluster 2: } \mu_2 = \frac{1}{3}[(5, 4) + (6, 5) + (7, 5)] = \left(\frac{18}{3}, \frac{14}{3}\right) = (6, 4.67)$$

## Convergence Check

No points change clusters with these new centroids  $\rightarrow$  Algorithm converges!

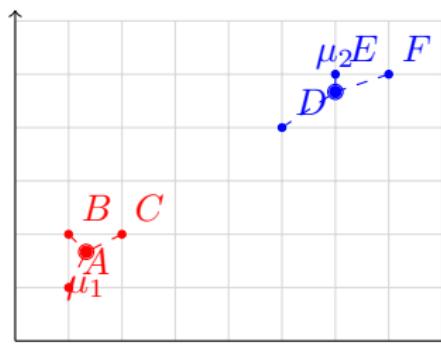


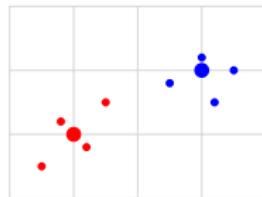
Figure: Final clustering after convergence

# Centroid Initialization Methods

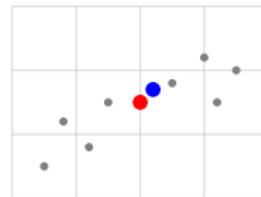
## The Problem

K-Means is sensitive to initial centroid positions → Can converge to local minima

Good Initialization



Bad Initialization



## Common Methods

- ① **Random initialization:** Choose  $k$  random points from dataset
- ② **K-Means++:** Smart initialization to spread out initial centroids
- ③ **Forgy method:** Choose  $k$  random data points as centroids
- ④ **MacQueen method:** Random partition then compute centroids

# K-Means++ Initialization

---

## Algorithm 2 K-Means++ Initialization

---

- 1: Choose first centroid  $\mu_1$  uniformly at random from data points
  - 2: **for**  $i = 2$  to  $k$  **do**
  - 3:   For each point  $x$ , compute  $D(x) = \text{distance to nearest centroid}$
  - 4:   Choose  $\mu_i$  with probability proportional to  $D(x)^2$
  - 5: **end for**
- 

In scikit-learn: `n_init` parameter

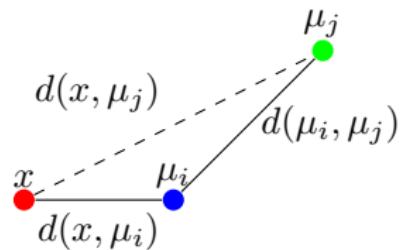
- `n_init`: Number of times algorithm is run with different centroid seeds
- Final results: Best output of `n_init` consecutive runs
- Default: `n_init=10`

# Accelerated K-Means (Elkan's Algorithm)

## Key Idea

Use triangle inequality to avoid unnecessary distance calculations

- Maintain lower bounds on distances
- If  $d(\mu_i, \mu_j) \geq 2d(x, \mu_i)$ , then  $x$  cannot be closer to  $\mu_j$
- Avoid computing  $d(x, \mu_j)$
- Speedup:** 2-10x faster for high-dimensional data



$$d(x, \mu_j) \geq |d(\mu_i, \mu_j) - d(x, \mu_i)|$$

# Mini-Batch K-Means

## Idea

Use random subsets (mini-batches) of data to update centroids

---

### Algorithm 3 Mini-Batch K-Means

---

- 1: Initialize centroids
- 2: **for** each iteration **do**
- 3:   Sample random mini-batch  $B$  from dataset
- 4:   **for** each point  $x$  in  $B$  **do**
- 5:     Assign  $x$  to nearest centroid
- 6:     Update centroid using moving average:

$$\mu_i \leftarrow \frac{n_i \mu_i + x}{n_i + 1}$$

where  $n_i$  is number of points assigned to centroid  $i$  so far

- 7:   **end for**
- 8: **end for**

# Finding Optimal Number of Clusters: Elbow Method

Inertia (Within-Cluster Sum of Squares)

$$\text{Inertia} = \sum_{i=1}^k \sum_{x \in C_i} \|x - \mu_i\|^2$$

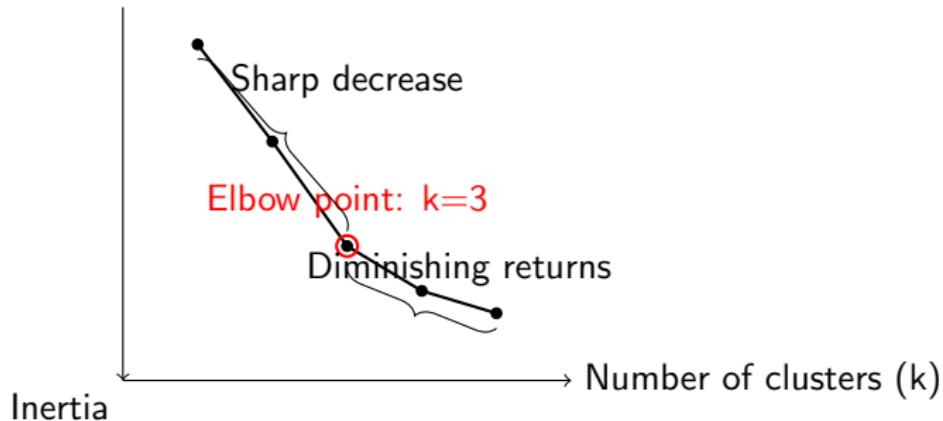


Figure: Elbow method: Choose k where inertia decrease slows down

# Silhouette Score Method

## Silhouette Score for a point

$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}}$$

where:

- $a(i)$ : Average distance to points in same cluster
- $b(i)$ : Average distance to points in next nearest cluster

## Interpretation

- $s(i) \in [-1, 1]$
- Close to 1: Well clustered
- Close to 0: Between clusters
- Close to -1: Wrong cluster

# Limitations of K-Means

## 1. Assumes Spherical Clusters

Does not work well for non-spherical clusters

## 2. Sensitive to Outliers

Computes Centroids based on all points even if they are outliers

## 3. Requires Specifying K

Must know or guess number of clusters

## 4. Scale Dependent

Features must be scaled (normalized)

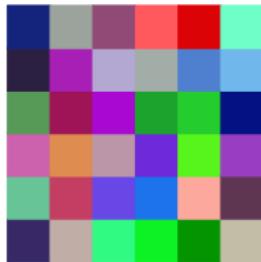
## 5. Equal Variance Assumption

Assumes clusters have similar sizes and densities

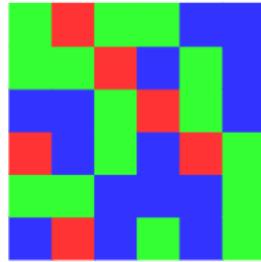
# Application: Image Segmentation with K-Means

## Idea

Treat each pixel as a data point in RGB space, cluster to find dominant colors



Original Image



Segmented ( $k=3$ )

## Process

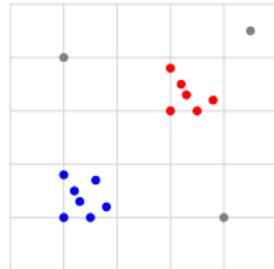
- ① Reshape image to pixels  $\times$  RGB values
- ② Apply K-Means with desired  $k$  (number of colors)
- ③ Replace each pixel with its cluster centroid color
- ④ Reshape back to image dimensions

# DBSCAN: Density-Based Spatial Clustering

## Key Idea

Cluster = dense region of points separated by low-density regions

- **Density-based:** Finds arbitrarily shaped clusters
- **Robust:** Handles outliers well
- **Parameters:**
  - $\epsilon$  (eps): Neighborhood radius
  - minPts: Minimum points to form dense region



## Core Concepts

- **Core point:** Has at least minPts points within  $\epsilon$  radius
- **Border point:** Within  $\epsilon$  of a core point but not core itself
- **Noise point:** Neither core nor border point

# DBSCAN: Understanding the Key Concepts

## 1. $\epsilon$ -Neighborhood (The "Friend Circle")

**In words:** All points that are within a distance  $\epsilon$  from point  $p$ . Think of it as drawing a circle of radius  $\epsilon$  around  $p$  - any point inside this circle is in its  $\epsilon$ -neighborhood.

**Mathematically:**

$$N_\epsilon(p) = \{q \in D \mid \text{distance}(p, q) \leq \epsilon\}$$

## 2. Core Point (The "Popular" Point)

**In words:** A point is a core point if it has at least  $\text{minPts}$  neighbors within its  $\epsilon$ -neighborhood (including itself). These are the points that form the dense centers of clusters.

**Mathematically:**

$$\text{Point } p \text{ is a CORE POINT if: } |N_\epsilon(p)| \geq \text{minPts}$$

# DBSCAN: How Points Connect to Form Clusters

## 3. Directly Density-Reachable (Immediate Connection)

**In words:** Point  $q$  is directly density-reachable from point  $p$  if:

- $p$  is a core point (has enough neighbors)
- $q$  is within  $p$ 's  $\epsilon$ -neighborhood (close to  $p$ )

Think:  $q$  is directly connected to the popular point  $p$ .

**Mathematically:**  $q$  is **directly density-reachable** from  $p$  if:

- ①  $p$  is a core point
- ②  $q \in N_\epsilon(p)$

## 4. Density-Reachable (Connected Through Friends)

**In words:** Point  $q$  is density-reachable from point  $p$  if you can travel from  $p$  to  $q$  through a chain of points where each point is directly density-reachable from the previous one. Like a friendship chain:  $p$  knows  $p_2$ ,  $p_2$  knows  $p_3$ , ..., eventually reaching  $q$ .

# DBSCAN: What Makes a Cluster?

## 5. Cluster Definition (The Complete Group)

**In words:** A cluster is a group of points where:

- **Completeness:** If a point  $p$  is in the cluster, then ALL points that are density-reachable from  $p$  must also be in the cluster.
- **Connectivity:** Any two points in the cluster are connected through density-reachability (directly or indirectly).

**Mathematically:** A cluster  $C$  satisfies:

- ① If  $p \in C$  and  $q$  is density-reachable from  $p$ , then  $q \in C$
- ② For any  $p, q \in C$ ,  $p$  and  $q$  are density-connected

### Key Insight

DBSCAN finds clusters by starting from a core point and "growing" the cluster by adding all density-reachable points. Points that aren't density-reachable from any core point are marked as noise/outliers.

# Toy Example: DBSCAN Algorithm

## Parameters

$\epsilon = 1.2$ , minPts = 3

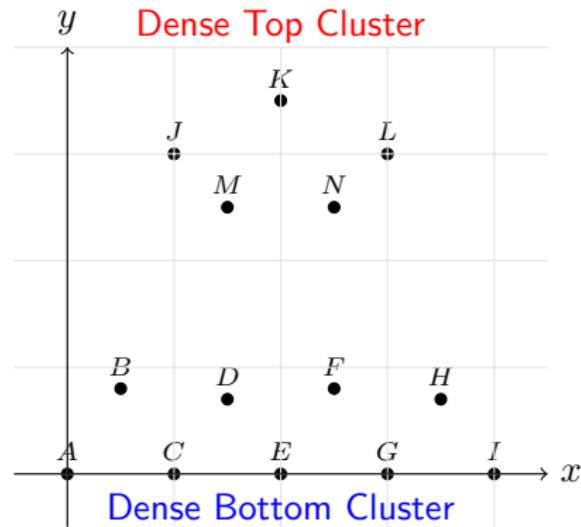
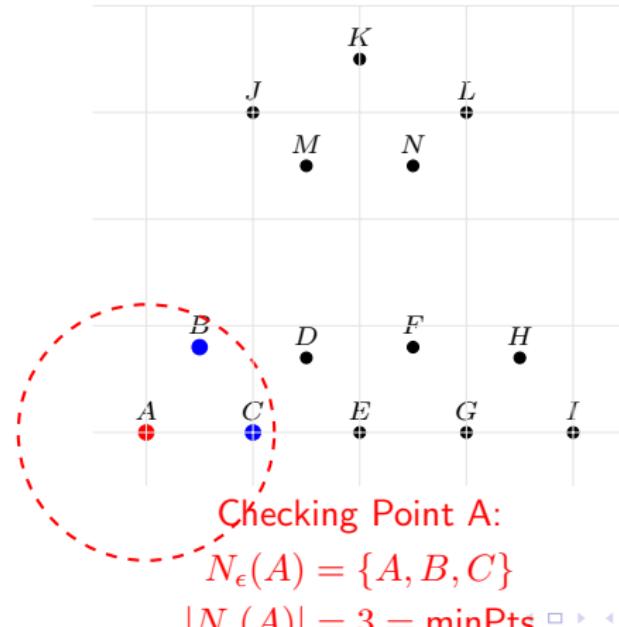


Figure: Two clearly separated clusters with  $\epsilon = 1.2$

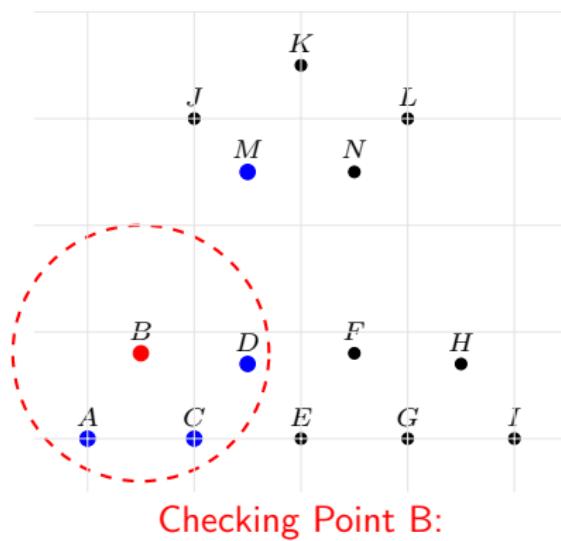
# Step 1: Identify Core Points - Check Point A

## Core Point Condition

A point is a **core point** if it has at least  $\text{minPts}=3$  points within  $\epsilon = 1.2$  radius (including itself)



## Step 2: Check Point B

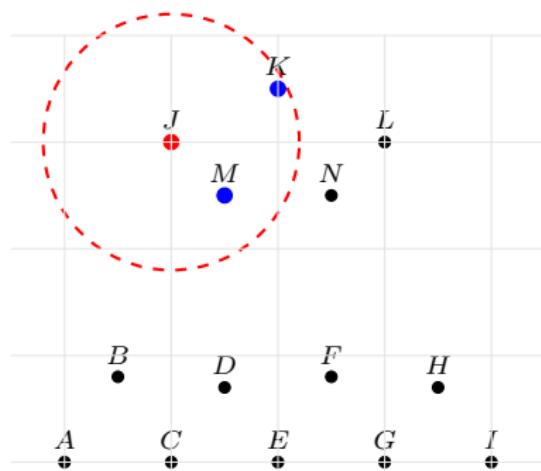


$$\text{Distance}(B, M) = \sqrt{(1.5 - 0.5)^2 + (2.5 - 0.8)^2} \approx 1.97 > 1.2$$

$$N_\epsilon(B) = \{A, B, C, D\}$$

$$|N_\epsilon(B)| = 4 \geq 3 \Rightarrow \text{CORE POINT}$$

# Step 3: Check Point J (First point in top cluster)



Checking Point J (Top Cluster):

$$\text{Distance}(J, K) = \sqrt{(2 - 1)^2 + (3.5 - 3)^2} \approx 1.12 < 1.2$$

$$\text{Distance}(J, M) = \sqrt{(1.5 - 1)^2 + (2.5 - 3)^2} \approx 0.71 < 1.2$$

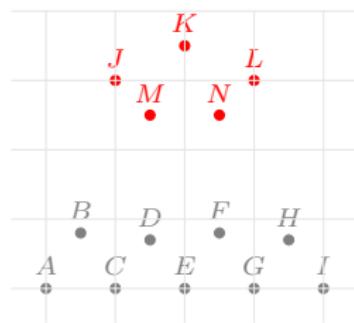
$$N_\epsilon(J) = \{J, M, K\} \Rightarrow |N_\epsilon(J)| = 3$$

$\Rightarrow J$  is a CORE POINT

# Step 4: All Core Points Identified

Core Point Analysis Complete

After checking all points (showing selected results):



Core Points

**Bottom Cluster:**

- Core: A, B, C, D, E, F, G, H
- Non-core: I (edge point)

**Top Cluster:**

- Core: J, K, L, M, N (all points)

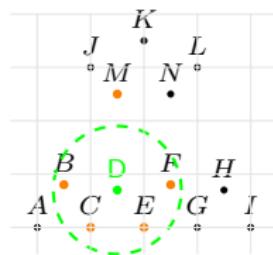
Why I is non-core

- $N_\epsilon(I) = \{H, I\}$  only
- $|N_\epsilon(I)| = 2 < 3$  (minPts)
- But I is within  $\epsilon$  of core point H
- $\Rightarrow I$  is a BORDER point

# Step 5: Start Clustering - Randomly Select Starting Point D

## DBSCAN Algorithm Begins

- All points unvisited initially
- Randomly select unvisited core point D
- Create Cluster 1
- Add D to Cluster 1
- Begin expansion from D



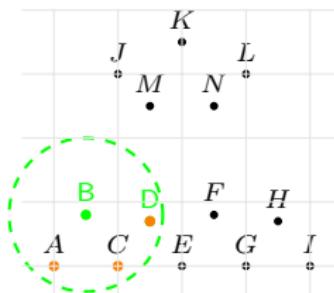
Starting with core point D  
Create Cluster 1

Neighbors within  $\epsilon = 1.2$ : B, C, E, F  
M is NOT within  $\epsilon$  (distance  $\approx 1.8 > 1.2$ )

# Step 6: Process B from Queue

## Expand Cluster 1

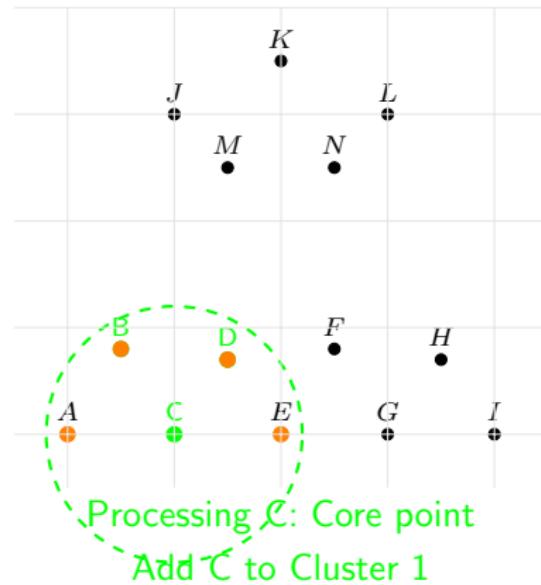
- Take B from queue
- B is core point (we identified earlier)
- Add B to Cluster 1
- Add B's unvisited neighbors to queue



Processing B: Core point  
Add B to Cluster 1

B's new neighbors: A, C (D already visited)  
Queue: C, E, F, A  
Cluster 1: {B, D}

## Step 7: Process C from Queue

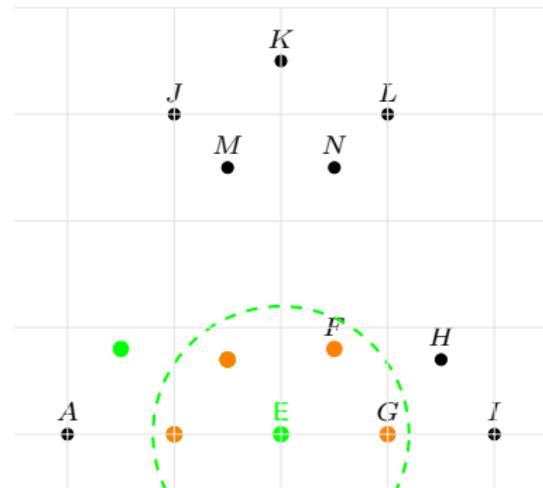


C's new neighbor: E (A already in queue)

Queue: E, F, A, E (E appears twice - will be deduplicated)

Cluster 1: {B, C, D}

## Step 8: Process E from Queue



Processing E: Core point  
Add E to Cluster 1

E's new neighbors: F, G (C, D already in cluster)

Queue: F, A, G

Cluster 1: {B, C, D, E}

## Step 9: Continue Expansion - Process F, A, G

### Processing Queue

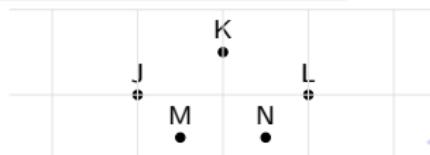
- **F**: Core point, adds to Cluster 1, adds H to queue
- **A**: Core point, adds to Cluster 1, no new neighbors
- **G**: Core point, adds to Cluster 1, adds H, I to queue



Next in queue

### Current State

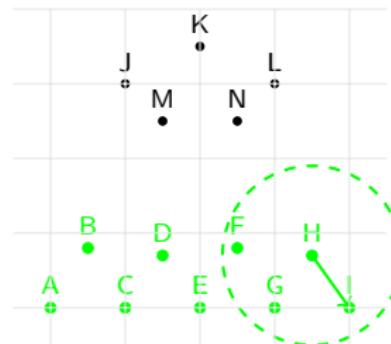
- Cluster 1: {A, B, C, D, E, F, G}
- Queue: H, I
- H is core point, I is border point



# Step 10: Process H and I - Complete Cluster 1

Final Points in Bottom Cluster

- **H:** Core point, adds to Cluster 1
- **I:** Border point (not core), but within  $\epsilon$  of H
- I is density-reachable from D through chain D-E-F-G-H-I



Processing H: Core point

Processing I: Border point (connected to H)

Cluster 1 COMPLETE! {A, B, C, D, E, F, G, H, I}

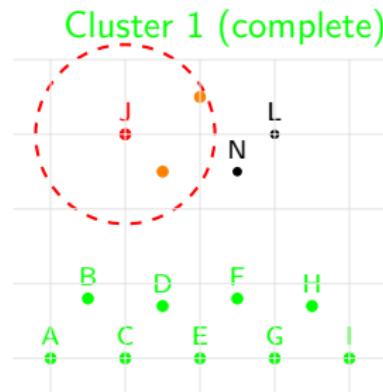
All bottom points clustered

Top cluster (J-N) still unvisited

# Step 11: Start Cluster 2 - Randomly Select J

Find Next Unvisited Point

- All points A-I are visited and in Cluster 1
- Next unvisited point: J (in top cluster)
- J is core point (we identified earlier)
- Create Cluster 2

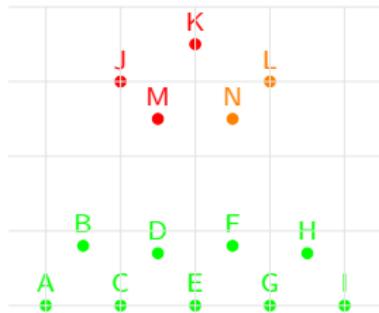


Start Cluster 2 with core point J  
 $N_{\epsilon}(J) = \{J, M, K\}$   
Neighbors to expand: M, K

# Step 12: Expand Cluster 2 - Process M and K

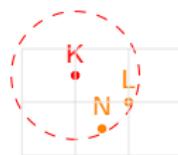
## Process Queue

- **M:** Core point, adds to Cluster 2
- M's neighbors: J (visited), K (in queue), N, L
- Add N, L to queue
- **K:** Core point, adds to Cluster 2
- K's neighbors: J, M (both visited), L, N
- L, N already in queue



Cluster 2: {J, M, K}  
Queue: L, N

All points in top cluster will connect

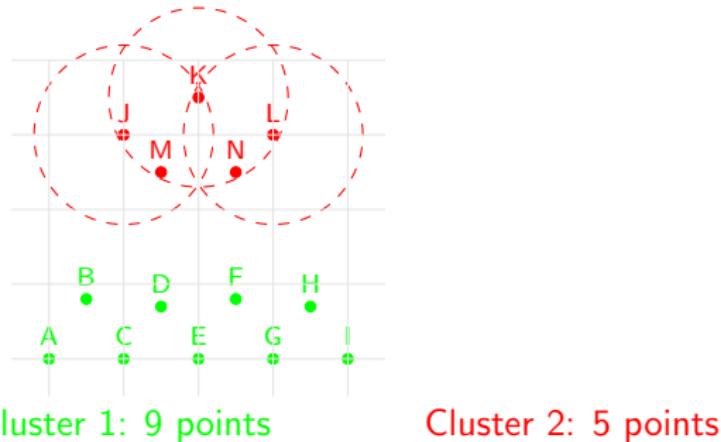


K connects to L and N

# Step 13: Complete Cluster 2

## Final Expansion

- **L**: Core point, adds to Cluster 2, connects to N
- **N**: Core point, adds to Cluster 2, connects to L, M, K
- All points in top cluster are density-connected



Cluster 1: 9 points

Cluster 2: 5 points

All points clustered - NO NOISE!  
DBSCAN successfully found 2 natural clusters

# DBSCAN Algorithm Summary

## What We Demonstrated

- ① **Core Point Identification:** Check each point's  $\epsilon$ -neighborhood
- ② **Cluster Growth:** Start from unvisited core point, expand via density-reachability
- ③ **Multiple Clusters:** Find all density-connected components
- ④ **Noise Detection:** Points not density-reachable from any core (none in our example)

## Key Insights from Our Example

- **Parameter Choice Matters:**  $\epsilon = 1.2$  perfectly separates the two clusters
- **Density-Based:** Finds clusters of arbitrary shape (both clusters are non-linear)
- **Robust:** Handles border points (like I) correctly
- **Order Independence:** Could start with any core point, would find same clusters

# DBSCAN: Advantages and Limitations

## Advantages

- **Arbitrary shapes:** Finds non-convex clusters
- **Robust to outliers:** Identifies noise points
- **No need for k:** Determines number of clusters automatically
- **Hierarchical:** Can find nested clusters with right parameters

## Limitations

- **Parameter sensitive:**  $\epsilon$  and minPts choice critical
- **Border points:** Can belong to multiple clusters (ambiguity)
- **Varying density:** Struggles with clusters of different densities
- **High-dimensional data:** Distance measures become less meaningful

## Choosing $\epsilon$ and minPts

- **minPts:** Rule of thumb:  $\text{minPts} \geq \text{dimensions} + 1$
- $\epsilon$ : Use k-distance plot (sorted distances to k-th nearest neighbor)
- Domain knowledge often needed

# K-Means vs DBSCAN: Comparison

Aspect	K-Means	DBSCAN
Cluster Shape	Spherical, convex	Arbitrary, non-convex
Outliers	Sensitive, affects centroids	Robust, identifies as noise
Parameters	Number of clusters (k)	$\epsilon$ (eps) and minPts
Number of Clusters	Must specify k	Determines automatically
Complexity	$O(n \cdot k \cdot d \cdot i)$	$O(n \log n)$ with spatial indexing
Scaling	Requires feature scaling	Distance-based, needs scaling
Use Case	Well-separated, spherical clusters	Density-based, arbitrary shapes

Table: Comparison of K-Means and DBSCAN

- **K-Means:** When you expect spherical clusters, know approximate k, need fast results
- **DBSCAN:** When clusters have arbitrary shapes, outliers exist, unknown k



# Summary

## Key Takeaways

- ① **K-Means:** Centroid-based, minimizes within-cluster variance, spherical clusters
- ② **DBSCAN:** Density-based, finds arbitrary shapes, robust to outliers
- ③ Both have strengths and weaknesses
- ④ Choice depends on data characteristics and problem requirements

## Practical Considerations

- Always preprocess data (scale features)
- Use visualization to understand cluster shapes
- Consider hybrid approaches
- Validate with domain knowledge

# References and Further Reading

-  MacQueen, J. (1967). Some methods for classification and analysis of multivariate observations.
-  Arthur, D., Vassilvitskii, S. (2007). k-means++: The advantages of careful seeding.
-  Elkan, C. (2003). Using the triangle inequality to accelerate k-means.
-  Ester, M., Kriegel, H., Sander, J., Xu, X. (1996). A density-based algorithm for discovering clusters in large spatial databases.
-  Pedregosa et al. (2011). Scikit-learn: Machine Learning in Python.

## Tools and Libraries

- scikit-learn: KMeans, DBSCAN, MiniBatchKMeans
- Visualization: Matplotlib, Seaborn, Plotly
- Large datasets: Spark MLlib, HDBSCAN